



**POLITECNICO**  
MILANO 1863

## **Hybrid Thread-MPI Parallelization for the ADR Equation**

Advection-Diffusion-Reaction Equation via FEM

**Peng Rao    Jiali Claudio Huang    Ruiying Jiao**

M.Sc. High Performance Computing — Politecnico di Milano

- 1 Problem Statement
- 2 FEM Discretization
- 3 Implementation
- 4 Parallelization Strategy
- 5 Experimental Results
- 6 Conclusions

## Strong Formulation

Find  $u : \Omega \rightarrow \mathbb{R}$  such that:

## ADR Equation with Mixed BCs

$$\begin{cases} -\nabla \cdot (\mu \nabla u) + \nabla \cdot (\beta u) + \gamma u = f & \text{in } \Omega \\ u = g & \text{on } \Gamma_D \\ \nabla u \cdot \mathbf{n} = h & \text{on } \Gamma_N \end{cases}$$

- $\Omega \subset \mathbb{R}^d$ ,  $d = 2, 3$
- $\mu > 0$ : diffusion coefficient
- $\beta \in [L^\infty(\Omega)]^d$ : advection field
- $\gamma \geq 0$ : reaction coefficient
- $f \in L^2(\Omega)$ : source term
- $\Gamma_D \cup \Gamma_N = \partial\Omega$

## Weak Formulation

Define  $V_0 = \{v \in H^1(\Omega) : v = 0 \text{ on } \Gamma_D\}$ .

## Variational Problem

Find  $u_0 \in V_0$  such that  $a(u_0, v) = F(v)$  for all  $v \in V_0$ , where:

$$a(u, v) = \int_{\Omega} \mu \nabla u \cdot \nabla v \, dx + \int_{\Omega} (\beta \cdot \nabla u) v \, dx + \int_{\Omega} \gamma u v \, dx$$

$$F(v) = \int_{\Omega} fv \, dx + \int_{\Gamma_N} \mu h v \, ds - a(u_g, v)$$

## Manufactured Solution

Domain:  $\Omega = [0, 1]^d$ ,  $d = 2, 3$ .

## Exact Solution

$$u_{\text{ex}}(\mathbf{x}) = \prod_{i=1}^d \sin(\pi x_i)$$

## Physical parameters:

- $\mu = 1.0$  (isotropic diffusion),  $\gamma = 0.1$  (reaction)
- Rotational advection:  $\beta = (-x_2, x_1, 0.1)^\top$  — divergence-free
- **Neumann BC** on right face  $x_1 = 1$ :  $h = -\pi \prod_{j \geq 2} \sin(\pi x_j)$
- **Dirichlet BC**  $u = 0$  on all other faces

**Triangulation**  $\mathcal{T}_h$ : hexahedral cells, mesh size  $h = \max_K \text{diam}(K)$ .  
**Space**  $V_h^k$ :  $Q_k$  Lagrangian elements (tensor-product degree  $k$ ).

### Discrete Problem

Find  $u_h \in V_{h,g}$  such that  $a(u_h, v_h) = L(v_h)$  for all  $v_h \in V_{h,0}$ .

Expanding in nodal basis  $\{\varphi_j\}$  gives the **linear system**:

$$\mathbf{A} \mathbf{U} = \mathbf{F}, \quad A_{ij} = \int_{\Omega} (\mu \nabla \varphi_j \cdot \nabla \varphi_i + (\boldsymbol{\beta} \cdot \nabla \varphi_j) \varphi_i + \gamma \varphi_j \varphi_i) \, dx$$

## A Priori Error Estimates

For  $u \in H^{k+1}(\Omega)$  and  $Q_k$  Lagrangian elements:

 **$H^1$  Seminorm**

$$|u - u_h|_{H^1} \leq Ch^k |u|_{H^{k+1}}$$

Rate:  $\mathcal{O}(h^k)$

 **$L^2$  Norm**

$$\|u - u_h\|_{L^2} \leq Ch^{k+1} |u|_{H^{k+1}}$$

Rate:  $\mathcal{O}(h^{k+1})$

- $H^1$  bound via **Céa's Lemma** + interpolation theory
- $L^2$  improvement via **Aubin-Nitsche** duality argument

**Assembly — WorkStream:**

$$A = \sum_{\text{cell}} P_{\text{cell}}^T A_{\text{cell}} P_{\text{cell}}$$

- **Worker:** local cell matrix  
(embarrassingly parallel)
- **Copier:** scatter to global  
(sequential)
- Sparse AIJ storage via **PETSc**

**Solver:**

- GMRES iterative solver
- **AMG** preconditioner  
(hypre)
- Memory:  
 $\mathcal{O}(N_{\text{dof}} \cdot (2k + 1)^d)$

Matrix-vector product computed *on the fly*:

$$\mathbf{v} = \sum_{e=1}^{N_{\text{el}}} P_e^T A_e (P_e \mathbf{u}), \quad (\text{no explicit } A \text{ stored})$$

### Optimizations:

- **Sum factorization:**  $\mathcal{O}(d(k+1)^{d+1})$  per cell
- **SIMD:** AVX over cell batches
- **Memory:**  $\mathcal{O}(N_{\text{dof}})$

### GMG Preconditioner:

- Chebyshev smoothers (deg. 5, range  $\times 15\text{--}20$ )
- Matrix-free transfer operators
- Coarse-grid Chebyshev solve

### MPI — Distributed Memory

- Mesh via p4est (Morton SFC)
- Each rank: subset of cells & DoFs
- Ghost cell exchange
- Parallel matrix/vector assembly
- Multigrid transfers across ranks

### TBB — Shared Memory

- Within each MPI rank
- MB: WorkStream parallel assembly
- MF: partition\_partition cell loop
- Parallel vector operations

### Key Finding

Pure MPI ( $28M \times 1T$ ) outperforms pure threading ( $1M \times 28T$ ) by **39×** at 28 cores.

# Why MPI Dominates

- **Memory bandwidth:** 28 threads share 1 controller; MPI uses multiple
- **NUMA effects:** remote-NUMA threads suffer  $3\text{--}4\times$  latency penalty; MPI ranks can be pinned locally
- **Cache efficiency:** threading shares 16.8M DoFs in one space → cache thrash; MPI distributes  $\approx 600\text{K}$  DoFs/rank → fits in LLC
- **Library optimization:** deal.II & PETSc are MPI-first; threading is secondary

## Recommendation

Use **pure MPI** (1 process/core). If hybrid is required, maximize MPI ranks and minimize threads per rank.

## Result

Both implementations confirm the expected  $L^2$ -error rate  $\mathcal{O}(h^3)$ .

## Theoretical:

$$\|u - u_h\|_{L^2} \leq Ch^{k+1} = \mathcal{O}(h^3)$$

## Observed:

- ✓ Matrix-based:  $\mathcal{O}(h^3)$
- ✓ Matrix-free:  $\mathcal{O}(h^3)$
- Both solvers are **correct**

- Céa's Lemma + Aubin-Nitsche

Time Complexity — Config: Hybrid  $7 \times 4$  (28 cores)

| $N_{\text{dof}}$ | AMG Matrix-Based |        |        | GMG Matrix-Free |         |        |
|------------------|------------------|--------|--------|-----------------|---------|--------|
|                  | Setup            | Assem. | Solve  | Setup           | Assem.  | Solve  |
| 1K               | 0.009s           | 0.001s | 0.009s | 0.011s          | 0.0001s | 0.013s |
| 17K              | 0.018s           | 0.010s | 0.025s | 0.021s          | 0.0004s | 0.044s |
| 263K             | 0.159s           | 0.047s | 0.508s | 0.189s          | 0.006s  | 0.556s |
| 4.2M             | 2.55s            | 1.61s  | 8.03s  | 2.26s           | 0.094s  | 6.98s  |
| 67M              | 40.7s            | 22.6s  | 155.5s | 37.1s           | 1.59s   | 125.3s |
| 268M             | 167.7s           | 92.7s  | 721.8s | 147.8s          | 6.10s   | 573.3s |

- **Crossover** at  $\approx 263K$  DoFs: MF faster beyond this point
- MF assembly:  $92.7s \rightarrow 6.1s$  (on-the-fly vs. explicit storage)
- GMG: **constant 6 iterations** vs. AMG 19–48 (growing with  $N$ )

# Memory Consumption

Table: Memory (MB) — Hybrid 7×4

| $N_{\text{dof}}$ | MB    | MF   |
|------------------|-------|------|
| 1K               | 0.37  | 0.03 |
| 17K              | 2.41  | 0.28 |
| 263K             | 28.2  | 4.1  |
| 4.2M             | 416   | 64   |
| 67M              | 6517  | 1026 |
| 268M             | 25979 | 4099 |

6.3× reduction at 268M DoFs

- Matrix-based: **26 GB**
- Matrix-free: **4.1 GB**

MB storage:  $\mathcal{O}(N \cdot (2k + 1)^d)$   
MF storage:  $\mathcal{O}(N)$  (vectors only)

*MF enables solving 6× larger problems on the same hardware.*

**At 28 cores, 16.8M DoFs:**

- MF total: 7.2s MB total: 13.5s
- MF speedup: **1.9×** over MB

**Parallel efficiency:**

- Matrix-free:  $\approx 74\%$  at 28 cores
- Matrix-based:  $\approx 50\%$  at 28 cores
- Larger problems → better efficiency

**MPI vs. Threading (28 cores):**

| Config | Time   | Eff. |
|--------|--------|------|
| 28M×1T | 7.2s   | 74%  |
| 14M×2T | 9.1s   | 58%  |
| 4M×7T  | 42.3s  | 10%  |
| 1M×28T | 281.8s | <2%  |

# Key Findings

1. **Correctness:** Both solvers achieve  $\mathcal{O}(h^{k+1})$   $L^2$  convergence — confirmed against manufactured solution.
2. **Memory (6.3×):** Matrix-free needs 4.1 GB vs. 26 GB at 268M DoFs.
3. **Algorithmic scalability:** GMG-MF maintains **constant 6 iterations** across  $10^3$ – $2.68 \times 10^8$  DoFs; AMG-MB grows from 19 to 48.
4. **Parallel efficiency:** MF achieves 74% at 28 cores vs. 50% for MB.
5. **MPI >> Threading:** Pure MPI outperforms pure threading by **39×** at 28 cores due to memory bandwidth, NUMA, and cache effects.

# Recommendations & Future Work

## Recommendations

- $N_{\text{dof}} > 250K$ : prefer **MF + GMG**
- Small problems / prototyping: MB is competitive
- Production: **pure MPI**, 1 process/core
- Memory-limited: MF solves  $6 \times$  larger problems

## Future Work

- 3D problems & higher  $k$
- GPU acceleration for MF operator
- Adaptive mesh refinement
- Time-dependent & nonlinear ADR

## Repository

<https://github.com/Peng-Rao/HybridADRSolver>